


# ARIMA Notebook

This is a brief report generated by R Markdown for the Bitcoin ARIMA prediction task.

The first step is to import data, which is sourced from Yahoo Finance.


```
library('ggplot2')
library('forecast')
library('tseries')
csv_path = '~/Downloads/BTC-USD (1).csv'
daily_data = read.csv(csv_path, header=TRUE, stringsAsFactors=FALSE)
daily_data$Date = as.Date(daily_data$Date, format="%d/%m/%Y")
ggplot(daily_data, aes(Date, Close)) + geom_line() + scale_x_date('month') + ylab("Daily Bitcoin Price")
xlab("")
```



The second step is to cleanse the data for outliers and missing values.

```
close_ts = ts(daily_data[, c('Close')])
daily_data$clean_close = tsclean(close_ts)
ggplot() +
  geom_line(data = daily_data, aes(x = Date, y = clean_close)) + ylab('Daily Bitcoin Price')
```

## Don't know how to automatically pick scale for object of type ts. Defaulting to continuous.



Next step is to define moving average.

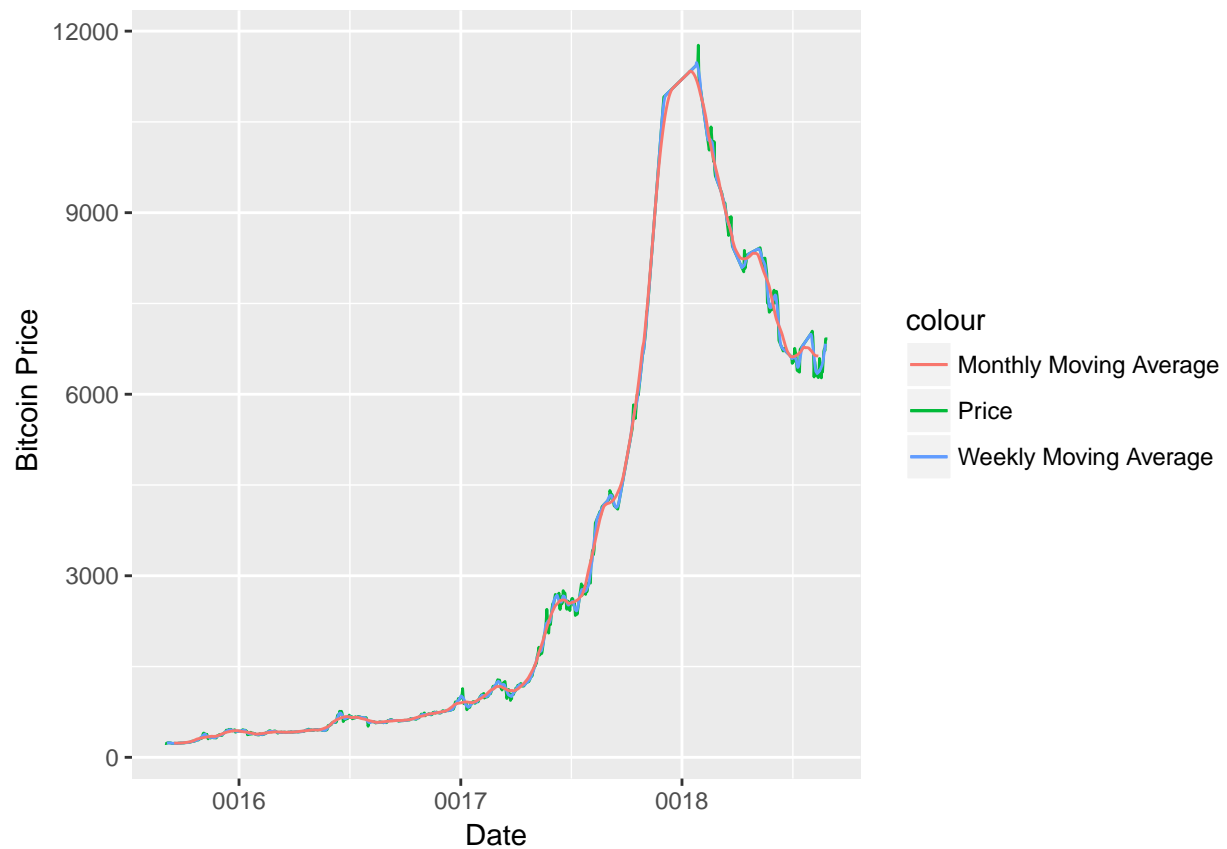
```
daily_data$close_ma = ma(daily_data$clean_close, order=7) # using the clean count with no outliers
daily_data$close_ma30 = ma(daily_data$clean_close, order=30)

ggplot() +
  geom_line(data = daily_data, aes(x = Date, y = clean_close, colour = "Price")) +
  geom_line(data = daily_data, aes(x = Date, y = close_ma, colour = "Weekly Moving Average")) +
  geom_line(data = daily_data, aes(x = Date, y = close_ma30, colour = "Monthly Moving Average")) +
  ylab(' Bitcoin Price')
```

## Don't know how to automatically pick scale for object of type ts. Defaulting to continuous.


## Warning: Removed 6 rows containing missing values (geom\_path).

## Warning: Removed 30 rows containing missing values (geom\_path).




Next step is data decomposition based on TS analysis.

```
close_ma = ts(na.omit(daily_data$close_ma), frequency=30)
decomp = stl(close_ma, s.window="periodic")
deseasonal_cls <- seasadj(decomp)
plot(decomp)
```




The next step is to make the model stationary, given the time dependency of the mean and variance. The two figures in the following help us to do that.


```
Acf(close_ma, main='')
```



```
Pacf(close_ma, main='')
```



```
close_d1 = diff(deseasonal_cls, differences = 2)
plot(close_d1)
```




```
adf.test(close_d1, alternative = "stationary")
```

```
## Warning in adf.test(close_d1, alternative = "stationary"): p-value smaller
## than printed p-value
##
## Augmented Dickey-Fuller Test
##
## data: close_d1
```

```
## Dickey-Fuller = -12.603, Lag order = 10, p-value = 0.01  
## alternative hypothesis: stationary
```


```
Acf(close_d1, main='ACF for Differenced Series')
```

### ACF for Differenced Series



```
Pacf(close_d1, main='PACF for Differenced Series')
```


## PACF for Differenced Series



The final step is to fit the model and generated the forecast. In the following, forecast is generated for the next 90 days to provide some visibility on the graph.

```
fit2 = arima(deseasonal_cls, order=c(1,1,9))  
tsdisplay(residuals(fit2), lag.max=15, main='Seasonal Model Residuals')
```

### Seasonal Model Residuals



```
fcast <- forecast(fit2, h=90)  
plot(fcast)
```

### Forecasts from ARIMA(1,1,9)

